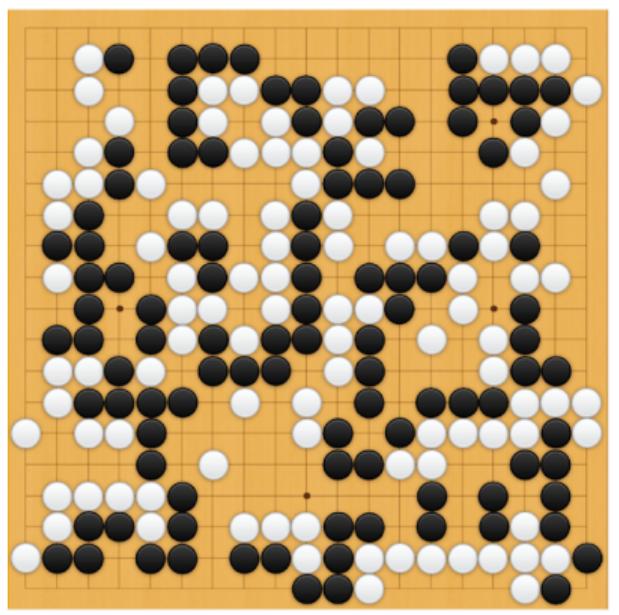


Statistical Inference in Reinforcement Learning

Chengchun Shi

Assistant Professor of Data Science
London School of Economics and Political Science

Developing AI with Reinforcement Learning



The image shows a Go board with black and white stones. On the right side, there is a banner with the text "THE ULTIMATE GO CHALLENGE GAME 3 OF 3" and the date "27 MAY 2017". Below the banner, there is a circular icon with a blue and white spiral pattern representing AlphaGo, followed by the text "AlphaGo" and "Winner of Match 3". Next to it is a circular portrait of a man with glasses, representing Ke Jie, with the text "Ke Jie" next to it. At the bottom, there is a large button with the text "RESULT B + Res".

THE ULTIMATE GO CHALLENGE
GAME 3 OF 3
27 MAY 2017

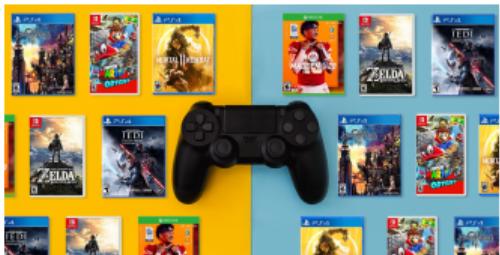
vs

AlphaGo
Winner of Match 3

Ke Jie

RESULT B + Res

Reinforcement Learning Applications



(a) Games



(b) Health Care



(c) Ridesharing



(d) Robotics



(e) Finance

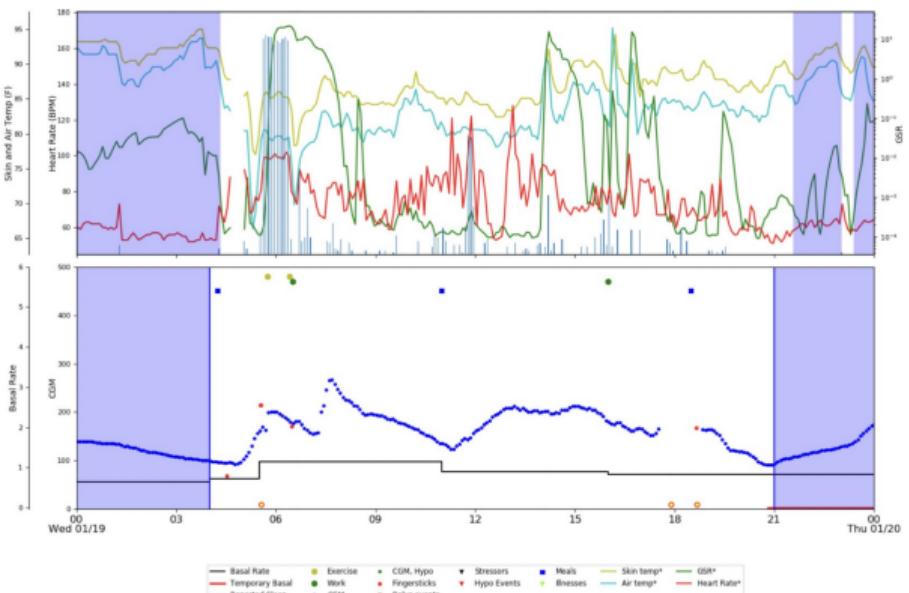


(f) Automated Driving

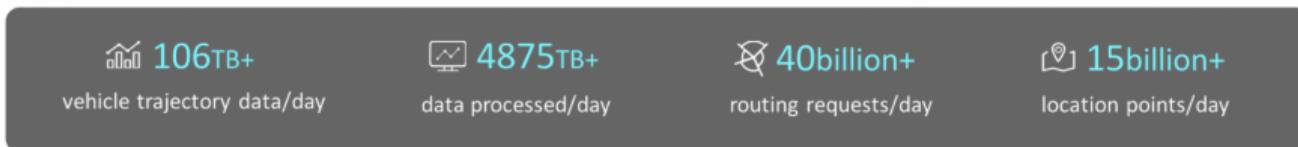
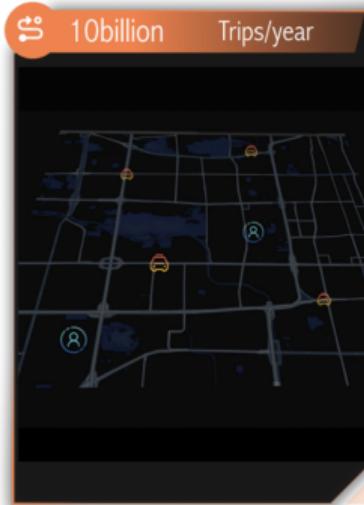
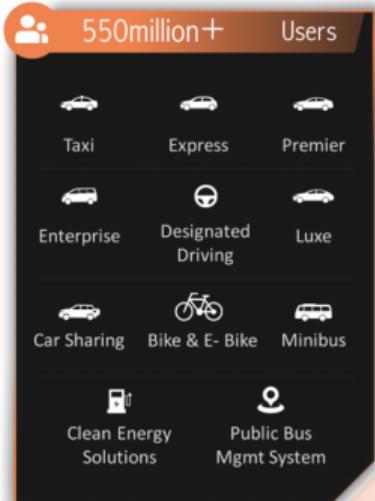
We focus on applications in **mobile health** (mHealth) and **ridesharing**

Applications in mHealth

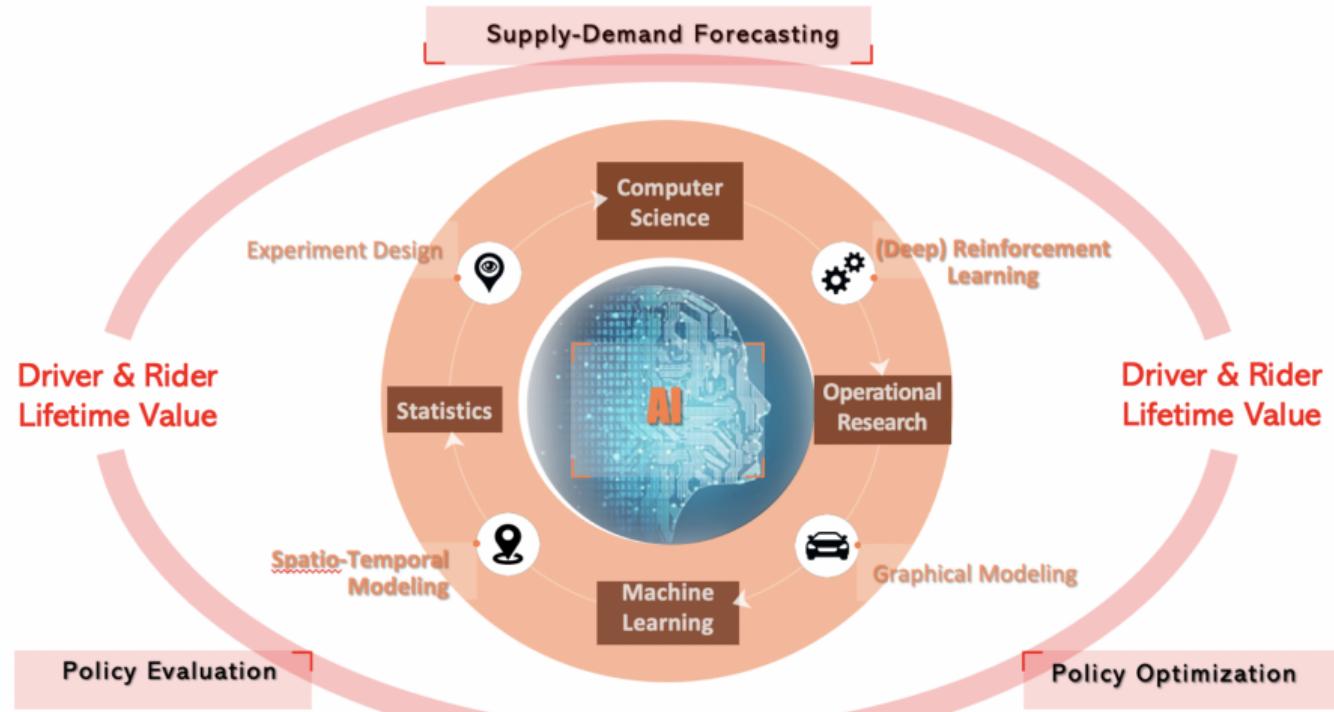
- Management of **Type-I diabetes**
- **Subject:** Patients with Type-I diabetes
- **Intervention:** Determine whether a patient needs to **inject insulin or not** based on their glucose levels, food intake, exercise intensity
- **Data:** OhioT1DM dataset (Marling and Bunescu, 2018)



Applications in Ridesharing



Applications in Ridesharing (Cont'd)



In this talk, we will focus on ...

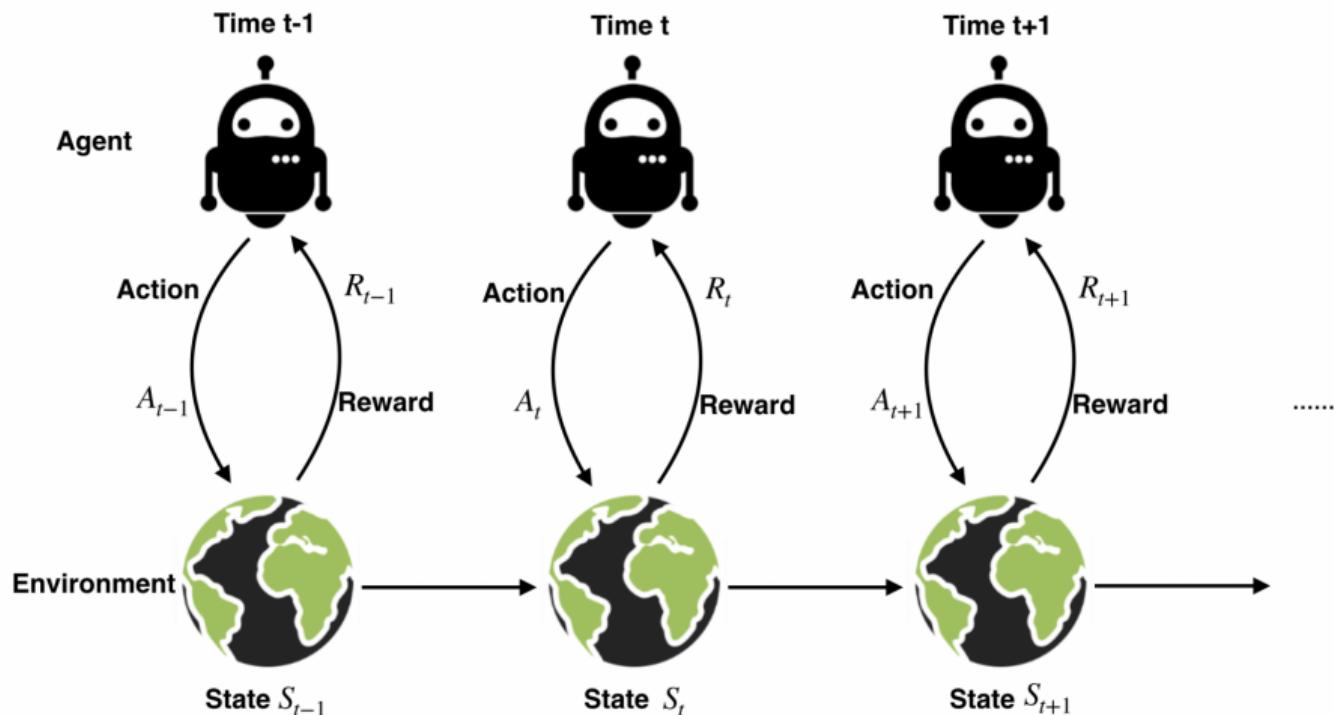
- **Statistical inference** in reinforcement learning (RL)
- Is statistical inference useful for RL?

Project I

Does the Markov Decision Process Fit the Data: Testing for the Markov Property in Sequential Decision Making

Joint work with Runzhe Wan, Wenbin Lu, Rui Song and Ling Leng
—ICML (2020)

Sequential Decision Making



Objective: find an optimal policy that maximizes the cumulative reward

The Agent's Policy

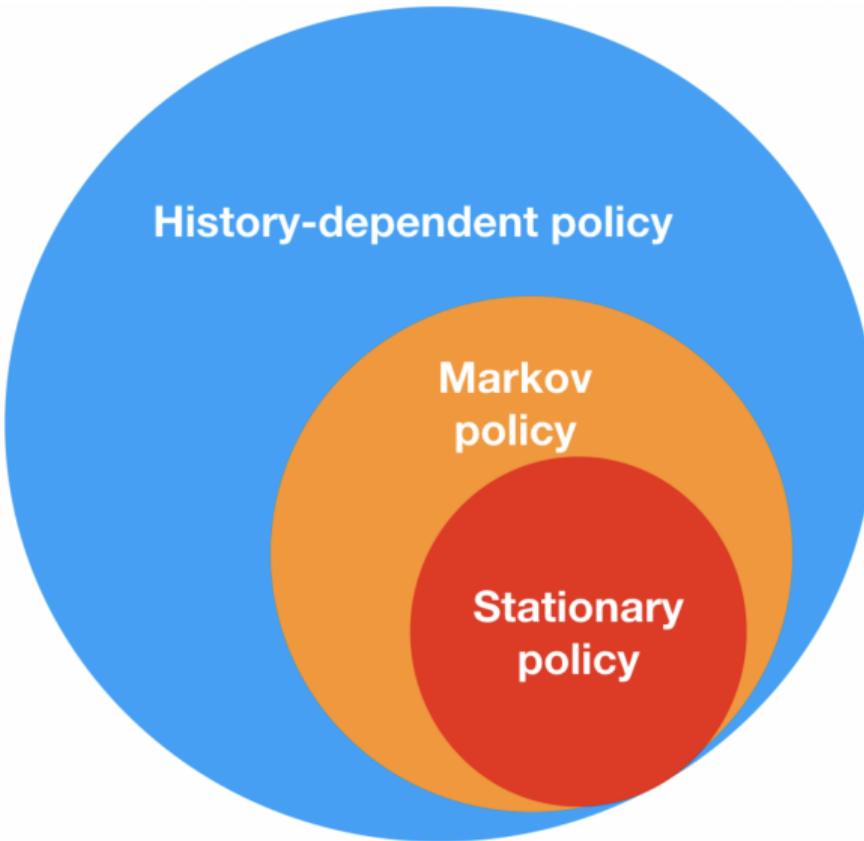
- The agent implements a **mapping** π_t from the observed data to a probability distribution over actions at each time step
- The collection of these mappings $\pi = \{\pi_t\}_t$ is called **the agent's policy**:

$$\pi_t(a|\bar{s}) = \Pr(A_t = a | \bar{S}_t = \bar{s}),$$

where $\bar{S}_t = (\mathcal{S}_t, \mathcal{R}_{t-1}, \mathcal{A}_{t-1}, \mathcal{S}_{t-1}, \dots, \mathcal{R}_0, \mathcal{A}_0, \mathcal{S}_0)$ is the set of **observed data history** up to time t .

- **History-Dependent Policy:** π_t depends on \bar{S}_t .
- **Markov Policy:** π_t depends on \bar{S}_t only through S_t .
- **Stationary Policy:** π is Markov & π_t is **homogeneous** in t , i.e., $\pi_0 = \pi_1 = \dots$.

The Agent's Policy (Cont'd)



Reinforcement Learning

- **RL algorithms:** trust region policy optimization (Schulman et al., 2015), deep Q-network (DQN, Mnih et al., 2015), asynchronous advantage actor-critic (Mnih et al., 2016), quantile regression DQN (Dabney et al., 2018).
- **Foundations** of RL:
 - **Markov decision process** (MDP, Puterman, 1994): ensures the optimal policy is *stationary*, and is *not* history-dependent.
 - **Markov assumption** (MA): conditional on the present, the future and the past are independent,

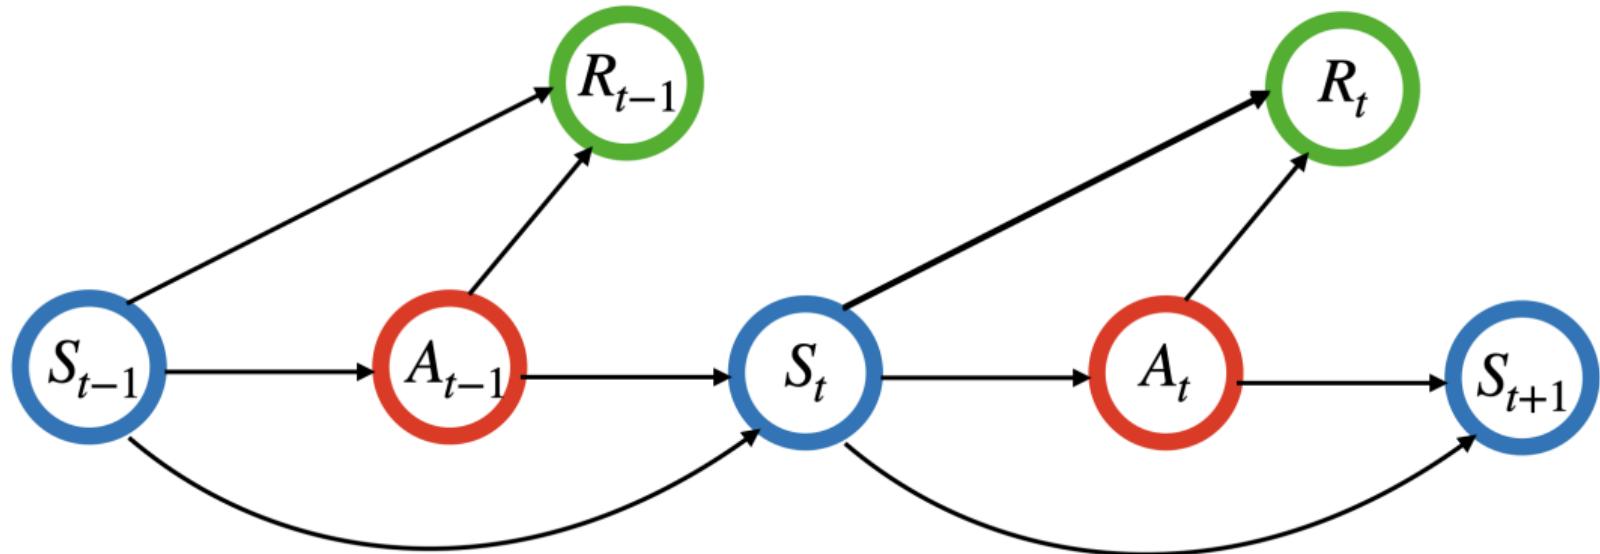
$$S_{t+1}, R_t \perp\!\!\!\perp \{(S_j, A_j, R_j)\}_{j < t} | S_t, A_t.$$

When R_t is a deterministic function of (S_t, A_t, S_{t+1})

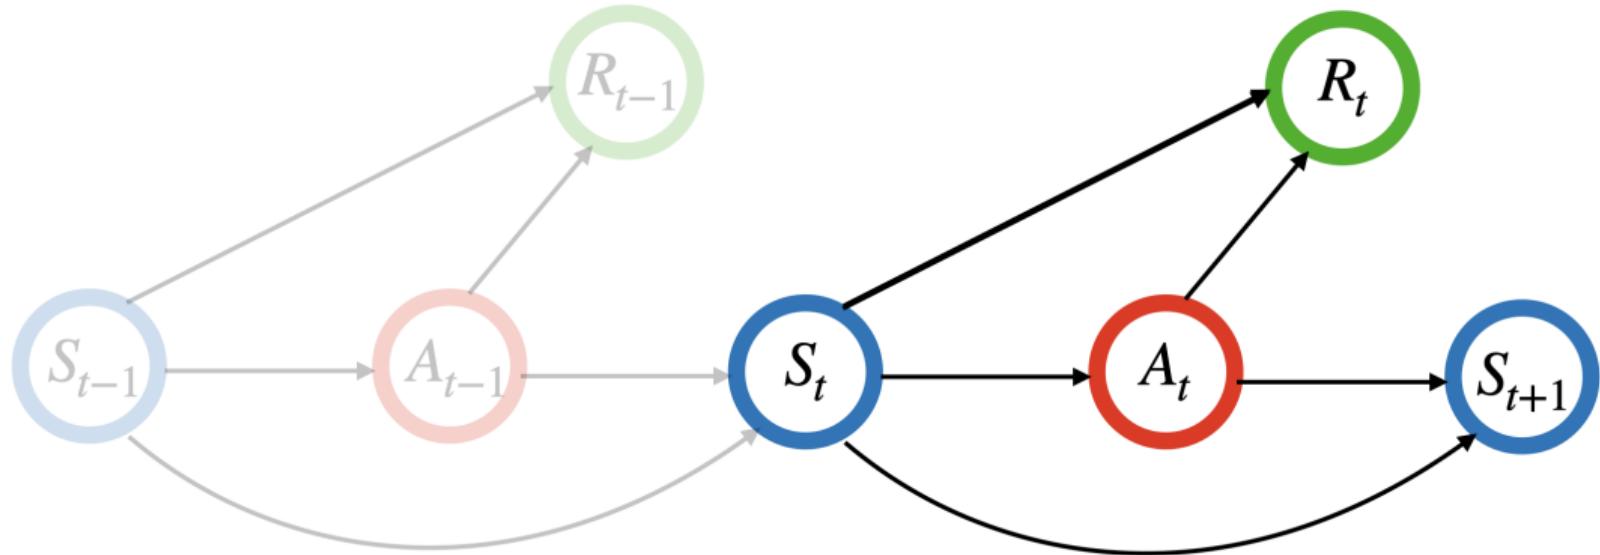
$$S_{t+1} \perp\!\!\!\perp \{(S_j, A_j)\}_{j < t} | S_t, A_t.$$

The Markov transition kernel is homogeneous in time

Markov Assumption



Markov Assumption



RL Models

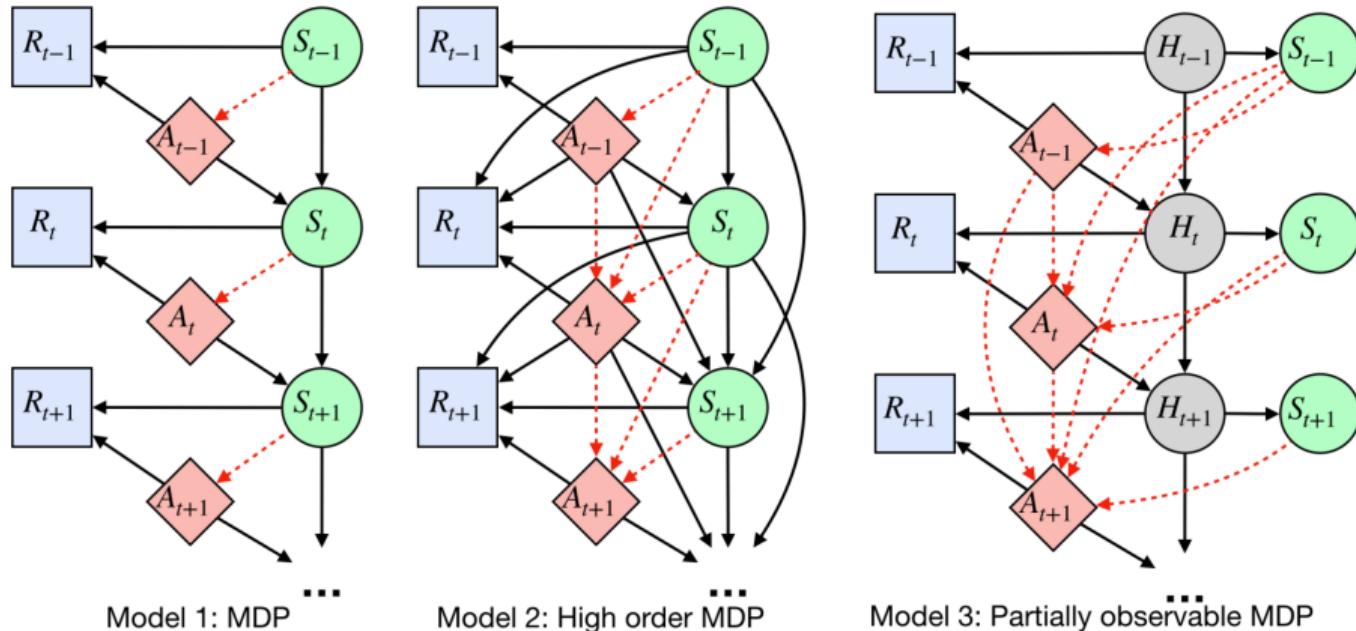


Figure: Causal diagrams for MDPs, HMDPs and POMDPs. The solid lines represent the causal relationships and the dashed lines indicate the information needed to implement the optimal policy. $\{H_t\}_t$ denotes latent variables.

Contributions

- **Methodologically**
 - propose a **forward-backward learning** procedure to test MA
 - **first** work on developing consistent tests for MA in RL
 - sequentially apply the proposed test for RL **model selection**
 - critical to **offline** domains:
 - For **under-fitted** models, any stationary policy is not optimal
 - For **over-fitted** models, the estimated policy might be very noisy due to the inclusion of many irrelevant lagged variables
- **Empirically**
 - identify the optimal policy in **high-order** MDPs
 - detect **partially observable** MDPs
- **Theoretically**
 - prove our test **controls type-I error** under a **bidirectional** asymptotic framework

Applications in High-Order MDPs

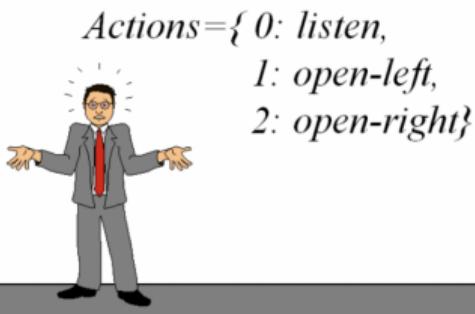
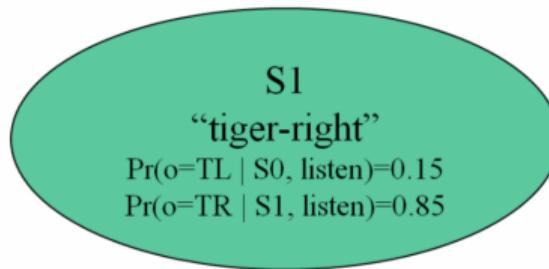
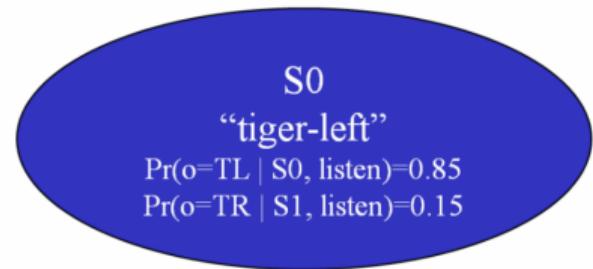
- **Data:** the OhioT1DM dataset (Marling & Bunescu, 2018)
- Measurements for 6 patients with type I diabetes over 8 weeks.
- One-hour interval as a time unit.
- **State:** patients' time-varying variables, e.g., glucose levels, food intake, exercise intensity
- **Action:** to inject insulin or not.
- **Reward:** the Index of Glycemic Control (Rodbard, 2009).

Applications in High-Order MDPs (Cont'd)

- **Analysis I:**
 - sequentially apply our test to determine the order of MDP
 - conclude it is a **fourth-order** MDP
- **Analysis II:**
 - split the data into training/testing samples
 - policy optimization based on **fitted-Q iteration**, by assuming it is a k -th order MDP for $k = 1, \dots, 10$
 - policy evaluation based on **fitted-Q evaluation**
 - use **random forest** to model the Q-function
 - repeat the above procedure to compute the average value of policies computed under each MDP model assumption

order	1	2	3	4	5	6	7	8	9	10
value	-90.8	-57.5	-63.8	-52.6	-56.2	-60.1	-63.7	-54.9	-65.1	-59.6

Applications in Partially Observable MDPs



Reward Function

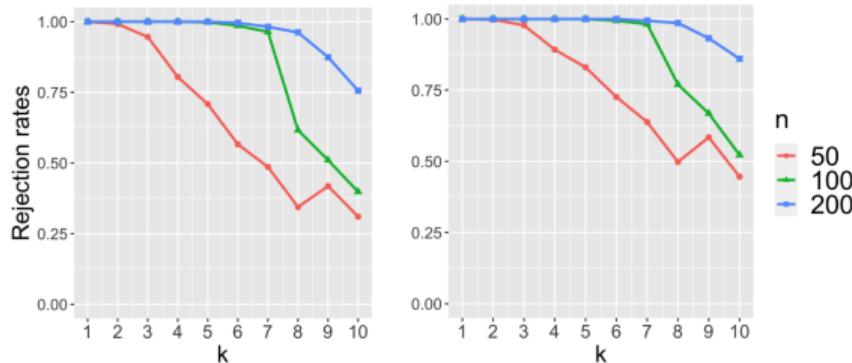
- Penalty for wrong opening: -100
- Reward for correct opening: +10
- Cost for listening action: -1

Observations

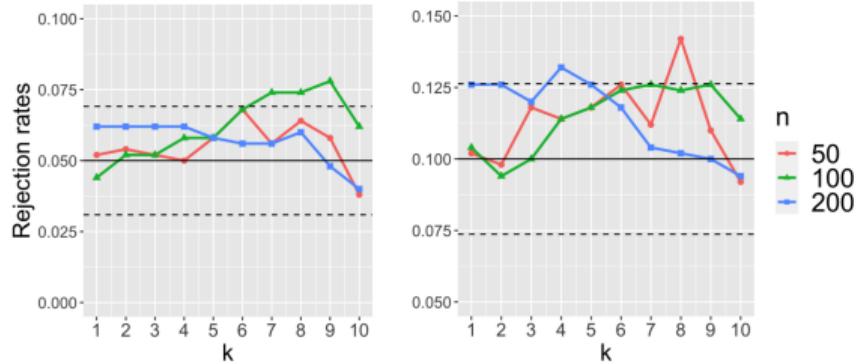
- to hear the tiger on the left (TL)
- to hear the tiger on the right (TR)

Applications in Partially Observable MDPs (Cont'd)

- Under the alternative hypothesis (MA is violated). $\alpha = (0.05, 0.1)$ from left to right.



- Under the null hypothesis (MA holds). $\alpha = (0.05, 0.1)$ from left to right.



Methodology

- **First** work to test MA in sequential decision making
- Existing approach in time series: Cheng and Hong (2012)
 - characterize MA based on the notion of **conditional characteristic function** (CCF)
 - use local polynomial regression to estimate CCF
- **Challenge:**
 - develop a valid test for MA in **moderate or high-dimensions**
 - the dimension of the state increases as we concatenate measurements over multiple time points in order to test for a high-order MDP.
- This motivates our **forward-backward learning** procedure.

Methodology (Cont'd)

Some key components of our algorithm:

- To deal with moderate or high-dimensional state space, employ modern machine learning (ML) algorithms to estimate CCF:
 - Learn CCF of S_{t+1} given A_t and S_t (**forward learner**)
 - Learn CCF of (S_t, A_t) given (S_{t+1}, A_{t+1}) (**backward learner**)
 - Develop a random forest-based algorithm to estimate CCF
 - Borrow ideas from the quantile random forest algorithm (Meinshausen, 2006) to facilitate the computation
- To alleviate the bias of ML algorithms, construct **doubly-robust** estimating equations by integrating forward and backward learners;
- To improve the power, construct a **maximum-type** test statistic;
- To control the type-I error, approximate the distribution of our test via **multiplier bootstrap** (Chernozhukov, et al., 2014).

Bidirectional Theory

- N the number of trajectories
- T the number of decision points per trajectory
- **bidirectional asymptotics:** a framework allows either N or $T \rightarrow \infty$
- large N , small T (Intern Health Study)



- small N , large T (OhioT1DM dataset)



- large N , large T (games)

Bidirectional Theory (Cont'd)

- (C1) Actions are generated by a fixed behavior policy.
- (C2) The observed data is exponentially β -mixing.
- (C3) The ℓ_2 prediction errors of forward and backward learners converge at a rate faster than $(NT)^{-1/4}$.

Theorem

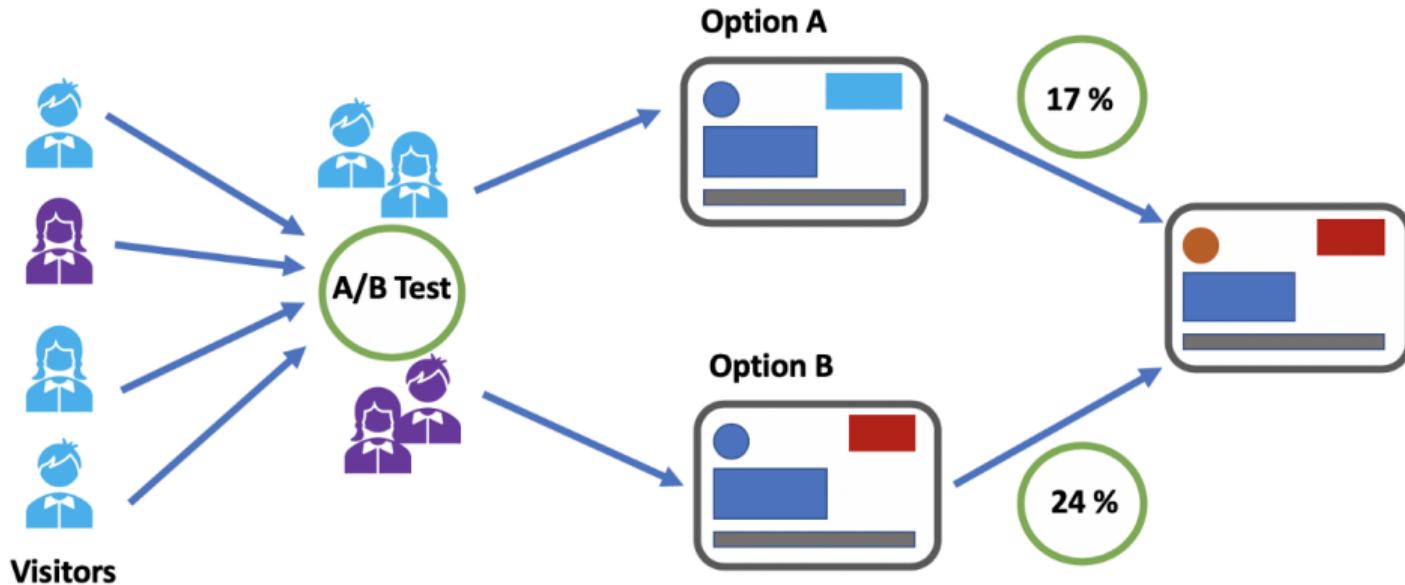
Assume (C1)-(C3) hold. Then under some other mild conditions, our test controls the type-I error asymptotically as either N or T diverges to ∞ .

Project II

Dynamic Causal Effects Evaluation in A/B Testing with a Reinforcement Learning Framework

*Joint work with Xiaoyu Wang, Shikai Luo, Hongtu Zhu, Jieping Ye and Rui Song
—JASA, accepted*

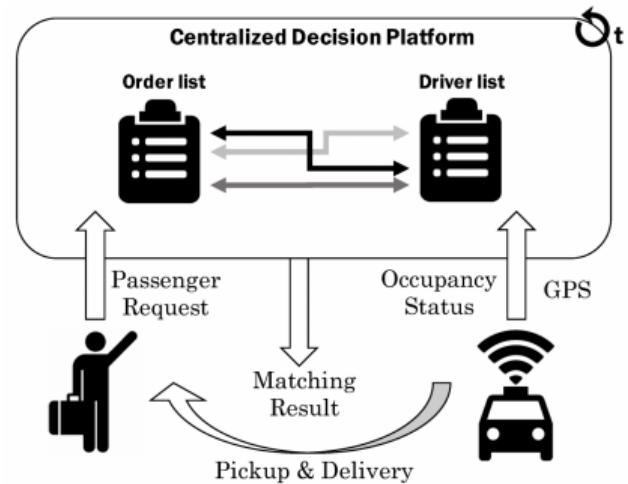
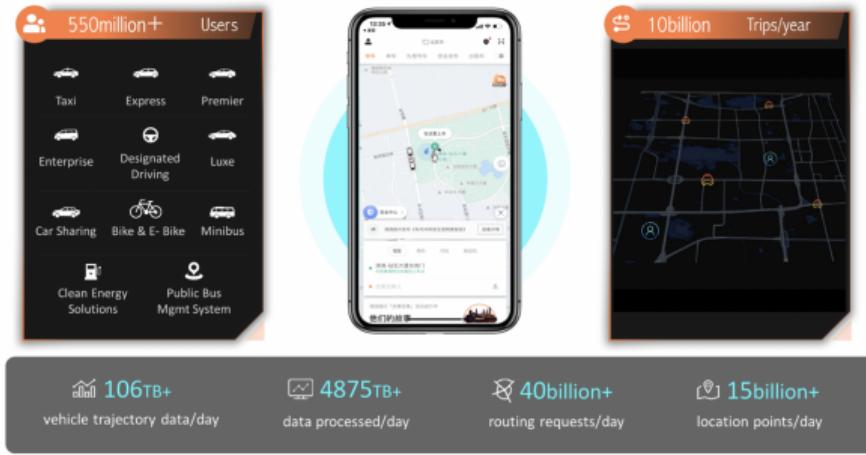
A/B Testing



Taken from

<https://towardsdatascience.com/how-to-conduct-a-b-testing-3076074a8458>

Motivation: Order Dispatch

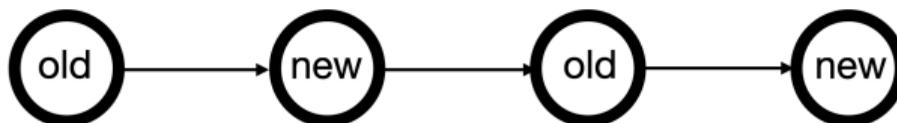


Our project is motivated by the need for comparing the **long-term rewards** of different **order dispatching** policies in **ridesharing platforms**

Challenges

1. The existence of **carryover effects**:

- Under the alternating-time-interval (or switchback) design



- Past actions will affect future outcomes

2. The need for **early termination**:

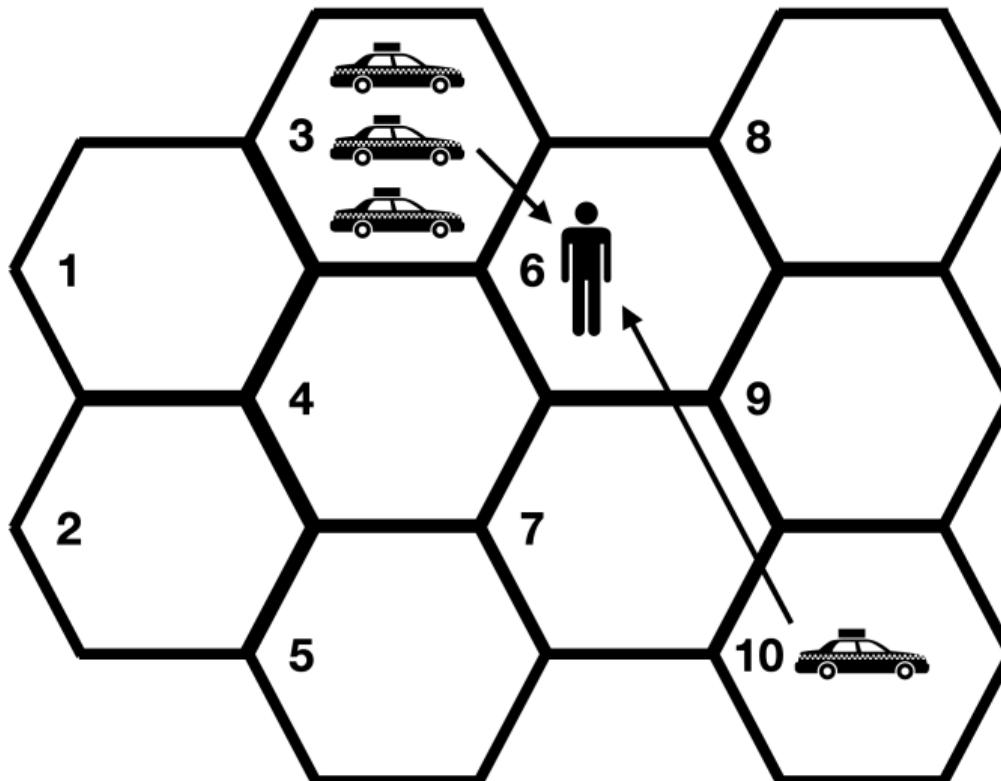
- Each experiment takes a considerable time (at most 2 weeks)
- Early termination to save time and budget

3. The need for **adaptive randomization**:

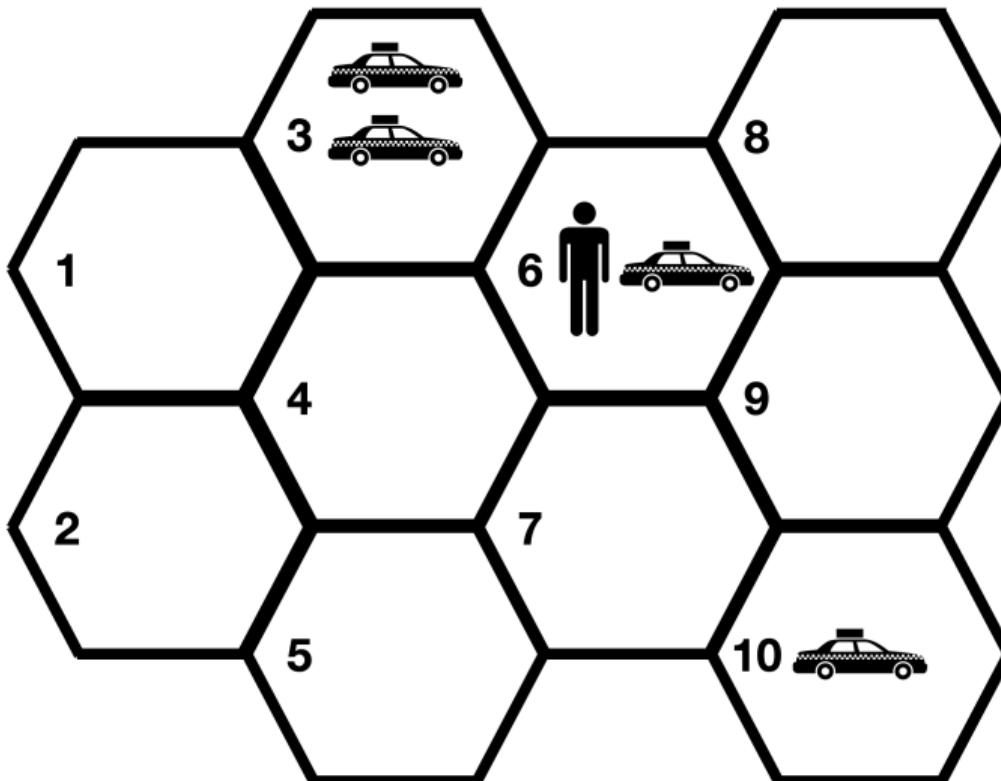
- Maximize the total reward (e.g., epsilon-greedy)
- Detect the alternative faster

To our knowledge, **no** existing test has addressed three challenges simultaneously

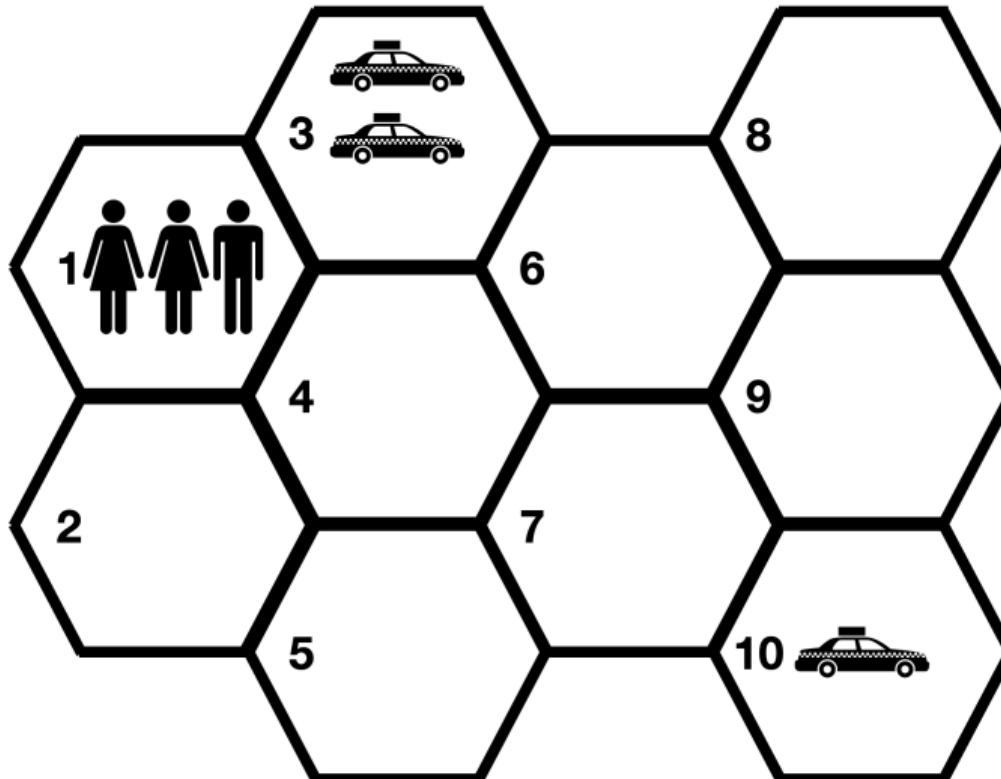
Illustration of the Carryover Effects



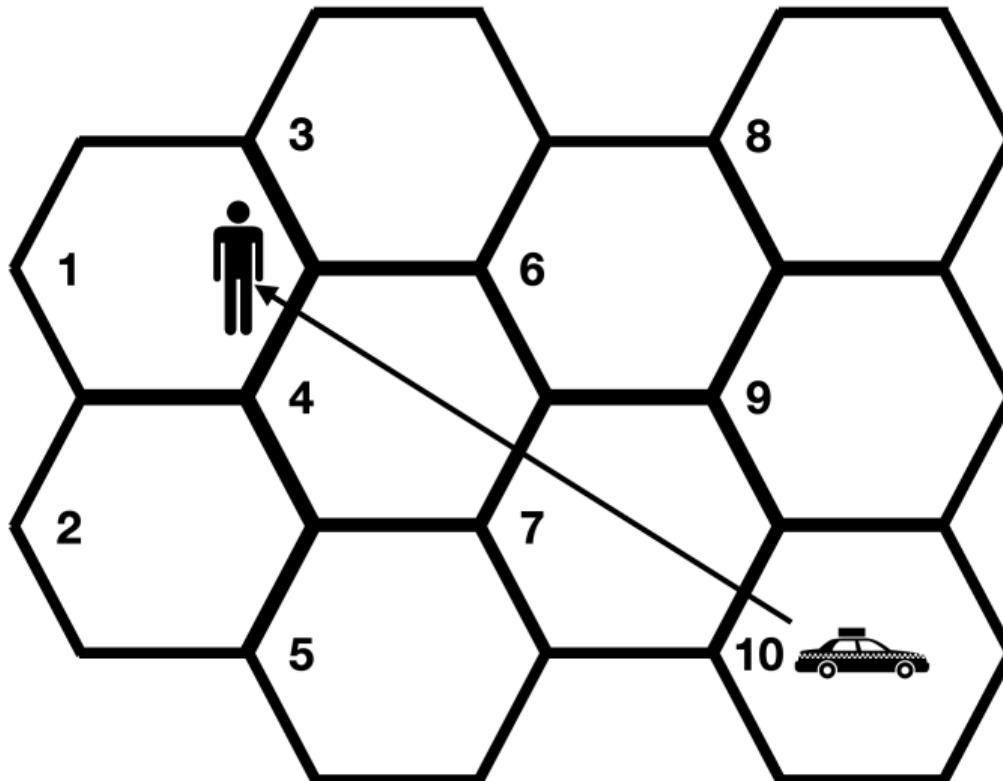
Adopting the Closest Driver Policy



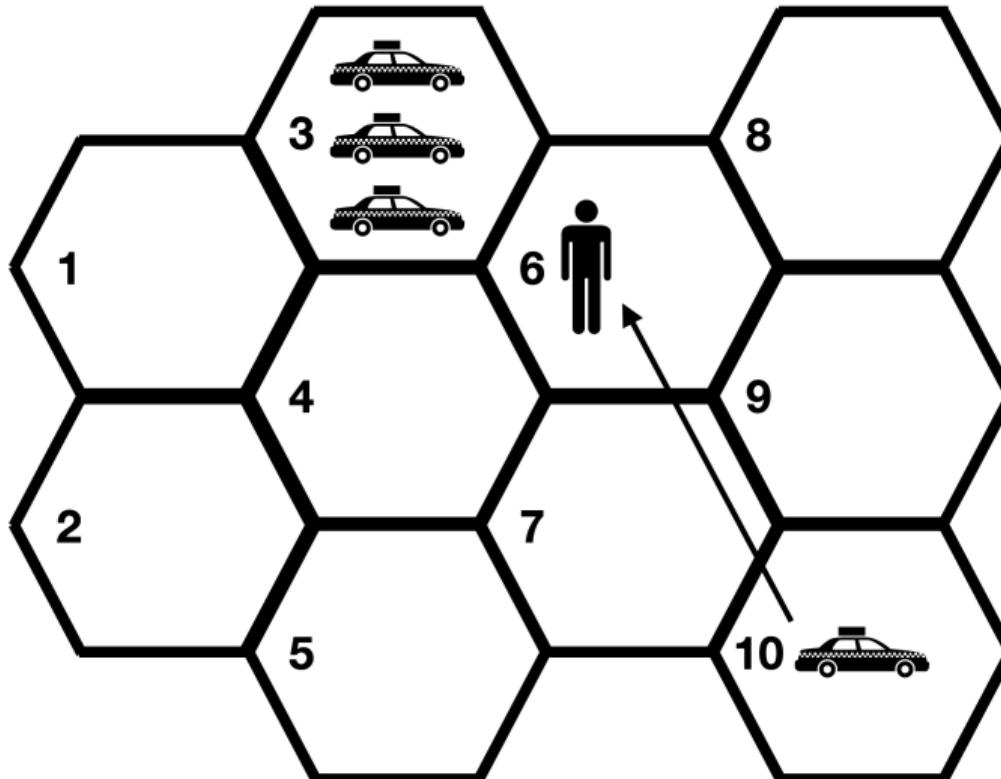
Some Time Later . . .



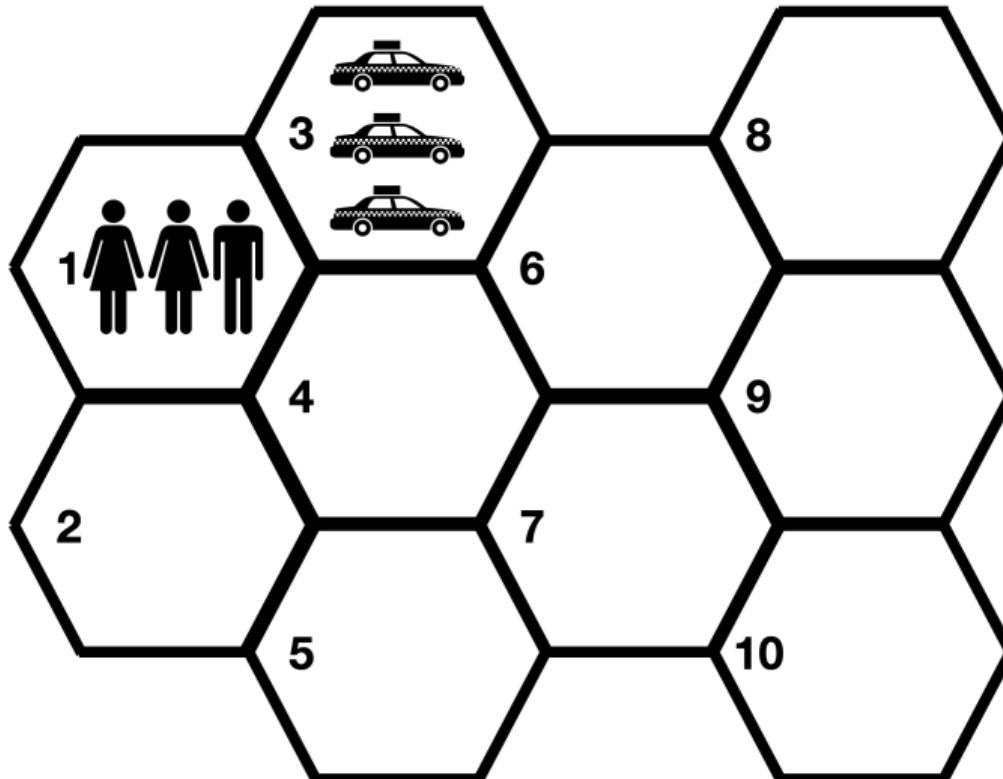
Miss One Order



Consider a Different Action



Able to Match All Orders



Existence of Carryover Effects

past actions → distribution of drivers → future rewards

Limitations of Existing A/B tests

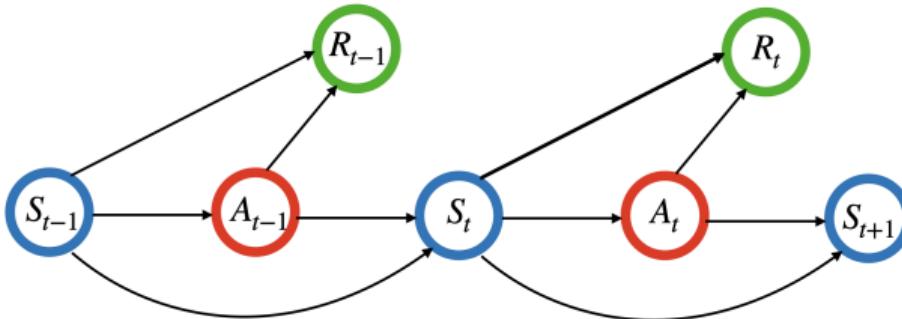
- Most existing tests **cannot** detect carryover effects
- \mathcal{H}_0 : The old policy ($A = 0$) has larger cumulative rewards
- \mathcal{H}_1 : The new policy ($A = 1$) has larger cumulative rewards
- **Example 1.** $S_t \sim N(0, 0.25)$, $R_t = S_t + \delta A_t$
- **Example 2.** $S_t = 0.5S_{t-1} + A_{t-1} + N(0, 0.25)$, $R_t = S_t$

Example 1	t-test 0.76	DML-based test 1.00	our test 0.98
Example 2	t-test 0.04	DML-based test 0.06	our test 0.73

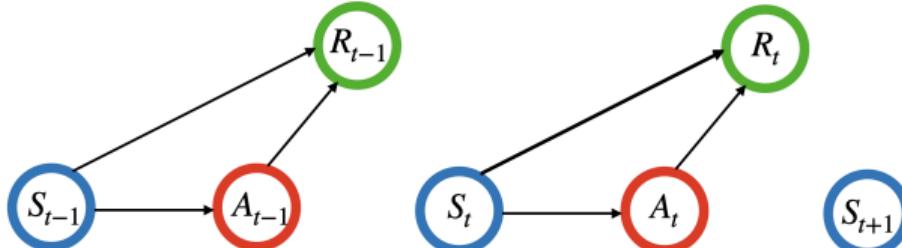
Table: Powers of t-test, DML-based test (Chernozhukov et al., 2018) and the proposed test with $T = 500, \delta = 0.1$

Contributions and Advances of Our Proposal

- Introduce an RL framework for A/B testing



1. A_{t-1} impacts R_t indirectly through its effect on S_t
 2. S_t shall include important **mediators** between A_{t-1} and R_t
- Most existing works require the independence assumption



Contributions and Advances (Cont'd)

Propose a test procedure for comparing long-term rewards of two policies

1. allows for **sequential monitoring**
2. allows for **online updating**
3. applicable to a wide range of designs, including the **Markov** design,
alternating-time-interval design and **adaptive** design

Methodology

- Apply **temporal difference learning** with **sieve** method to evaluate value difference and provide **uncertainty quantification** (Shi et al., 2021, JRSSB)
- Adopt the **α -spending approach** (Lan & DeMets, 1983) for sequential monitoring
- Develop a **bootstrap-assisted procedure** for determining the stopping boundary
 - The numerical integration method designed for classical sequential tests is **not** applicable in adaptive design, due to the carryover effects

Theory

Theorem (Validity and Consistency)

Under the Markov, alternating-time-interval or adaptive design, the proposed test can control type-I error and is consistent against alternatives that converge to the null at the parametric rate

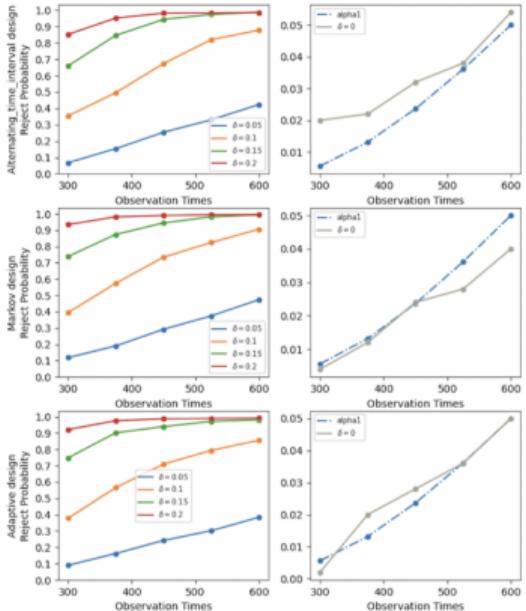
Theory (Cont'd)

Theorem (Undersmoothing and Efficiency)

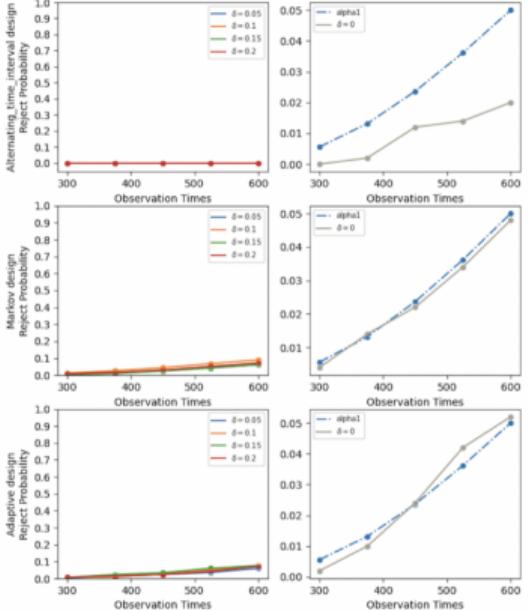
Suppose **sieve** method is used for function approximation in temporal difference learning.

1. **Undersmoothing** is not needed to guarantee that the resulting value estimator has a tractable limiting distribution.
 2. The value estimator is **semiparametrically efficient**.
- Sieve estimators of conditional expectations are **idempotent** (Shen et al., 1997)
 - The proposed test will **not** be overly sensitive to the number of basis functions
 - **Cross-validation** can be employed to select the basis functions

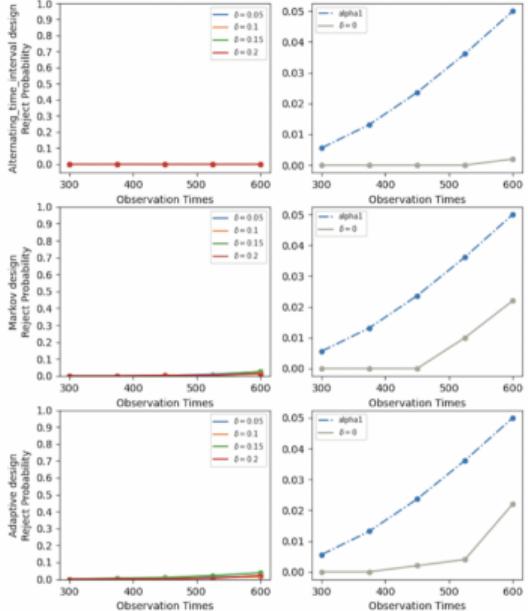
Simulation



(a) Power and size of our test



(b) Power and size of t test



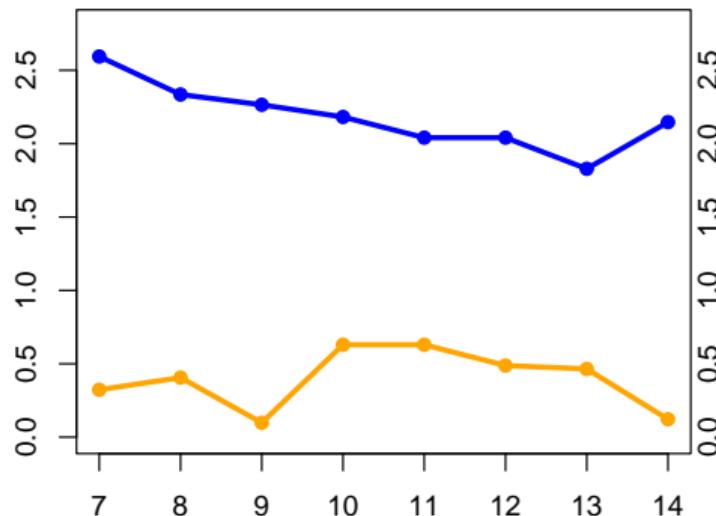
(c) Power and size of a version of the O'Brien Fleming sequential test

Application to Ridesharing Platform

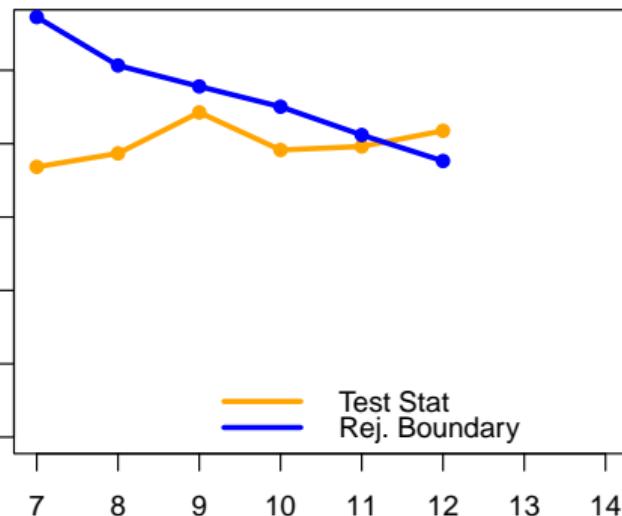
- **Data:** a given city from December 3rd to 16th (two weeks)
- **30 minutes** as one time unit, sample size = **672**
- **State:**
 1. number of drivers (supply)
 2. number of requests (demand)
 3. supply and demand equilibrium metric (mediator)
- **Action:** new policy **$A = 1$** v.s. old **$A = 0$**
- **Reward:** drivers' income
- The new policy is expected to have **better** performance

Application to Ridesharing Platform (Cont'd)

- The proposed test



(a) AA Experiment: Day



(b) AB Experiment: Day

- t-test: **fail** to reject \mathcal{H}_0 in A/B experiment with p-value 0.18

Project III

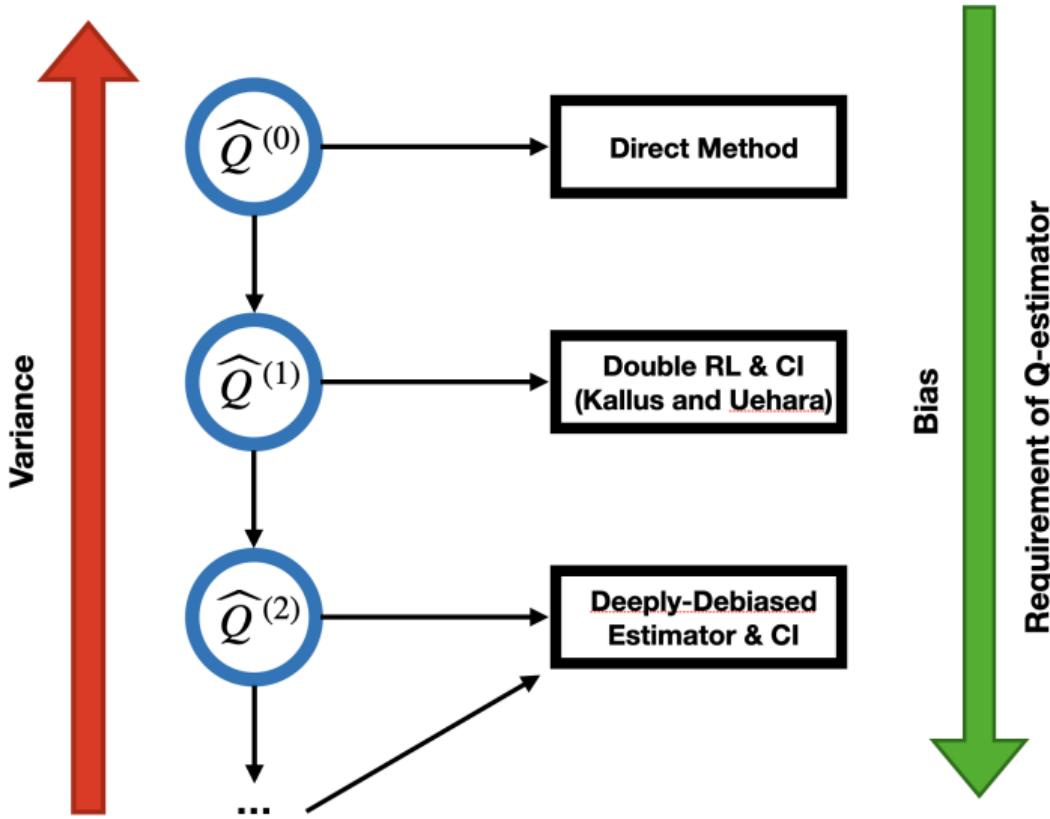
Deeply-Debiased Off-Policy Interval Estimation

*joint work with Runzhe Wan, Victor Chernozhukov, and Rui Song
—ICML, 2021 (long talk, top 3% of submissions)*

Off-Policy Interval Estimation

- **Objective:** Evaluate the impact of a target policy **offline** using historical data generated from a different behavior policy and provide rigorous **uncertainty quantification** (healthcare, automated driving, ridesharing, robotics, e.g.)
- Consider the reinforcement learning (e.g., MDP) setting
- Most existing methods focus on providing point estimators
- **Main idea:** Develop a **deeply-debiasing** process using higher order influence function (Robins et al., 2017)

Method



Thank You!

😊 Papers and softwares can be found on my personal website

callmespring.github.io